


Operating Principles

The rules behind the allocation model — N+1, fairness between tenants, conservative allocation, requests and limits, replicas, HA, day 2 operations and what it all costs — grouped by what the platform team owns, enforces, requires, runs and charges. On-prem clusters.



At a glance

One rule per topic. The box is for reviews: tick what a cluster or project meets today.

A. Capacity — what we can promise

	A1	N+1 (node failure tolerance)	Σ project quota still fits after the largest node — or failure domain — is gone.	☐
	A2	Conservative allocation	Sell at most 80 % in prod. The rest is reserve, not stock.	☐
	A3	Autoscaling and bursts	Quota covers the maximum (maxReplicas × requests), not the average.	☐

B. Tenancy — how we share it

	B1	Fairness between tenants	Every project gets what it pays for — quota is the only fairness Kubernetes has.	☐
	B2	Priority and preemption	Platform first, then prod. Priority never bypasses quota.	☐
	B3	Isolation beyond CPU and memory	Quota covers CPU and memory only; guard the rest with limits, policies and node pools.	☐

C. Workload standards — what we require from teams

	C1	Requests vs. limits (req < lmt)	Requests are the contract. Prod memory limit = request. CPU limits optional.	☐
	C2	Replicas	At least 2 replicas in prod, each able to carry the load alone.	☐
	C3	Workload HA	Spread over nodes, a PDB that allows one disruption, real probes.	☐

D. Platform reliability — what we run

	D1	Platform HA	3 control plane nodes, 3-node Rancher, 2+ ingress, node VMs on separate hosts.	☐
---	-----------	--------------------	--	--------------------------

E. Day 2 operations — keeping it true over time

	E1	Maintenance	Drains succeed unattended; upgrades go one node at a time.	☐
	E2	Monitoring and alerts	Alert before tenants notice: N+1, quota at 90 %, Pending pods, OOMKills.	☐
	E3	Reviews and changes	Monthly capacity check, quarterly project review, every quota change via Git.	☐

F. Finance — what it costs

	F1	Cost model and unit rates	Projects pay for allocated quota at published rates; the reserve is priced in, not billed.	☐
	F2	Rancher licence (CPU cores)	Rancher is licensed per CPU core: every core counts, sold or not; memory is licence-free.	☐
	F3	Cluster size: minimum worker nodes and N+1	Prod: at least 3 worker nodes; from 5 equal nodes on, N+1 fits inside the 20 % reserve.	☐



A. Capacity

what we can promise · How much can we sell, and still keep the promise?

Platform team owns. The guarantee in every other category depends on this one being right.

PLATFORM TEAM

Owens: sizes clusters, approves quota

TEAMS

Request quota with evidence



A1 N+1 (node failure tolerance)

Σ project quota still fits after the largest node — or failure domain — is gone.

Example: 4 nodes × 16 CPU



Allocatable 64 CPU



Platform 4


Σ project quota $\leq$ allocatable – largest node(s) – platform requests

- A cluster must be able to lose its **largest node** and still run every project's full quota.
- Check: Σ project quota $\leq$ allocatable – largest node(s) – platform requests. The planner and each cluster's dashboard show this as **Room for projects (N+1)** (allocation model, planner guide).
- Check it on **real node sizes**, not a flat percentage: with 3 nodes, one node is 33 % of the cluster; with 10, it's 10 %.
- Uneven node sizes: the largest node sets the reserve, so one oversized node makes the whole cluster expensive. Prefer nodes of equal size within a cluster.
- Prod: N+1 is mandatory. Acc/test: N+1 recommended. Dev: N+0 accepted (pods may stay Pending during a failure).
- N+2 only for clusters where two nodes can go down together (e.g. maintenance on one node plus a failure on another, or two VMs on one hypervisor host).
- A failure domain is whatever goes down together**: a bare-metal server, a hypervisor host running several node VMs, a rack, a power feed. N+1 counts failure domains, not Kubernetes nodes (open question Q17, D1).
- N+1 also covers **planned** outages: a node drained for patching is a failed node for capacity purposes. Without N+1, every upgrade is a capacity incident.
- Control plane / etcd nodes are not counted as capacity for projects.
- When a cluster breaks N+1 (red tile), it's a capacity trigger: add a node or stop approving quota (process, action D5 in 07).
- Minimum size**: 3 worker nodes in prod, better 5+ of equal size, so the N+1 reserve stays within the 20 % we keep anyway (F3).



A2 Conservative allocation

Sell at most 80 % in prod. The rest is reserve, not stock.



- Never sell more than the cluster can guarantee** in prod: Σ quota $\leq$ min(N+1 limit, 80 % of allocatable minus platform).

- The remaining 20 % is **not spare capacity to sell**: it absorbs rolling-update surge, rescheduling after a failure, platform growth and emergencies.
- Allocate on **requests**, never on limits or on "peak memory seen once".
- **Memory is allocated more carefully than CPU**: CPU shortage slows pods down (throttling), memory shortage kills them (OOMKill, eviction).
- Start new projects at their **measured or estimated need**, not their budget ceiling. Unspent budget is fine; unused quota blocks others.
- Increase quota in steps, on evidence: usage and requests trending towards the quota, Pending pods, a planned launch.
- Keep a small **unallocated platform reserve** per cluster for emergency increases, so an incident doesn't have to wait for new hardware (emergency increase).
- Overcommit ($\Sigma \text{ quota} > \text{capacity}$) only in **dev**, knowingly, and never promise dev capacity as guaranteed.
- Quota that is allocated but idle for two quarterly reviews is a candidate to reclaim (E3).
- Conservative $\neq$ stingy: a project that is short of quota will work around it (bigger pods elsewhere, fewer replicas, skipped HA). Say no to capacity we don't have, not to capacity that is justified.



A3 Autoscaling and bursts

Quota covers the maximum ($\text{maxReplicas} \times \text{requests}$), not the average.



- **HPA**: quota covers $\text{maxReplicas} \times \text{requests}$, not the average. $\text{minReplicas} \geq 2$ in prod.
- An HPA that hits its quota stops scaling silently (pods stay Pending with a quota error): alert on it (E2).
- **VPA** in recommendation mode only, as input for right-sizing; auto mode changes requests and therefore quota consumption without review.
- **No node autoscaling**: capacity is bought ahead, which is why allocation is conservative.
- **CronJobs and batch**: their requests count against quota while they run; schedule large jobs outside peak hours or give them their own quota.



B. Tenancy

how we share it · How do tenants share a cluster without hurting each other?

Platform team sets the rules and enforces them; tenants can't opt out.

PLATFORM TEAM

Sets the rules and enforces them

TEAMS

Stay within their Project



B1 Fairness between tenants

Every project gets what it pays for — quota is the only fairness Kubernetes has.

Tenant = Rancher Project



CPU on a busy node: shared by requests

requests

A: 4 CPU

B: 2 CPU

CPU time when all want more

$\frac{2}{3}$

$\frac{1}{3}$

Idle CPU may be used by anyone;
requested CPU is never taken away.

- The tenant is the **Rancher Project** (one per stream), not the namespace. Quota is set per Project per cluster; the project splits it over its namespaces itself (allocation model).
- Fairness = you get what you pay for:** a project's quota is guaranteed only because Σ quota fits the cluster (A1, A2). Without that, the scheduler is first come, first served.
- Kubernetes itself has **no fair-share scheduling** between tenants. Quota is the only mechanism that stops one tenant from taking all the capacity; that's why every project namespace gets a quota.
- CPU under contention** is shared in proportion to requests: a project with twice the requests gets twice the CPU time. Using idle CPU above requests is allowed (no CPU limits, ADR-0003) and never takes away another pod's requested CPU.
- Memory is not shared fairly:** a pod using more than its request is the first to be evicted under node pressure. Hence memory limit = request in prod (C1): nobody lives on borrowed memory.
- Same rules, same price for every project: one set of standards, published unit rates, same review cadence. Exceptions are written down with a reason and an end date.
- Special hardware (GPU, large-memory nodes) goes in **separate node pools** with taints, allocated and priced separately, so general workloads can't occupy it.
- Platform components (ingress, monitoring, logging, `cattle-*`, `kube-system`) are outside project quota and deducted from capacity before anything is sold — tenants don't pay for them via quota, and don't compete with them.
- Visibility is part of fairness: every project sees its own quota, requests and usage (dashboard, Kibana), and the showback compares projects on the same basis (ADR-0002).



B2 Priority and preemption

Platform first, then prod. Priority never bypasses quota.

- PriorityClasses, from high to low: **platform** (`system-cluster-critical` / `system-node-critical`), **prod**, **default**, **batch** / **best effort**.
- Preemption only makes sense when the cluster is short; with correct N+1 and headroom it should be rare.
- A high priority is **not a way around quota**: limit its use per namespace with a quota `scopeSelector` on `PriorityClass`, or a policy (Q5).
- Low-priority batch jobs can use capacity that is idle now, and are the first to be preempted when prod needs it.
- In mixed clusters (prod + non-prod on the same nodes, Q2), non-prod runs at lower priority.



B3 Isolation beyond CPU and memory

Quota covers CPU and memory only; guard the rest with limits, policies and node pools.

- Quota protects CPU and memory. **Not isolated** by default: disk I/O, network bandwidth, ephemeral storage, PIDs, API server load. These are where noisy neighbours come from.
- **Ephemeral storage:** set `requests/limits.ephemeral-storage` in standards, or a pod filling `/var/log` or `emptyDir` can get a node evicted.
- **Object counts** as guardrails, not billed: pods, services, `services.loadbalancers`, PVCs, ConfigMaps/Secrets (allocation model).
- **Persistent storage:** `requests.storage` quota per storage class if storage is significant (Q8).
- **Network:** NetworkPolicies per project (default deny between Projects; Rancher Project Network Isolation where available). Security isolation, not bandwidth.
- **Access:** project owners manage their own namespaces and quota split; they can't change other Projects or cluster-scoped objects (Q9).
- Heavy I/O or network workloads that disturb neighbours get dedicated nodes (taints + tolerations), allocated like any other node pool.



C. Workload standards

what we require from teams · What must a workload look like to get the guarantees?

Platform team defines the standards, sets defaults and checks compliance; teams implement them. A workload that doesn't follow them doesn't get the guarantees of A and B.

PLATFORM TEAM

Defines and checks standards, provides defaults

TEAMS

Implement them in manifests



C1 Requests vs. limits (req < lmt)

Requests are the contract. Prod memory limit = request. CPU limits optional.

Memory, prod

request

limit = request

Memory, non-prod

request

limit ≤ 2× request → eviction risk

CPU

request

no limit: may burst into idle

Quota and cost count requests. Memory above request is borrowed — and evicted first.

- **Requests are the contract:** what the scheduler reserves, what quota counts, what the project pays for.
- Requests must be set on every container, including sidecars and init containers (standards S1, S8).
- **Memory, prod: limit = request.** No borrowed memory; a pod is killed by its own limit (OOMKill, visible to the team), not evicted because a neighbour ran the node out of memory.
- **Memory, non-prod: limit ≤ 2× request** allowed. The gap between request and limit is overcommit: fine while pods don't all peak at once; the cost is evictions when they do.
- **Why req < lmt on memory is risky:** the node is filled by requests, but pods may grow to their limits. When the sum of actual usage exceeds the node, the kubelet evicts pods above their requests — often not the pod that caused it.
- **CPU: limits optional,** ≥ 2× request if set (ADR-0003). req < lmt on CPU is safe: CPU is compressible, and requests still guarantee each pod its share.
- Tight CPU limits (limit = request) cause throttling even on an idle node — bad for latency and JVM startup.
- **QoS:** with no CPU limit, pods are **Burstable**, not **Guaranteed**. Acceptable; the eviction risk comes from memory above request, which prod rules out anyway.
- **BestEffort** (no requests) is not allowed anywhere: it can't be accounted for and is evicted first (S5).
- Requests should be **right-sized** to ~P95 of usage (S6): too high wastes quota, too low causes throttling and evictions.
- Runtimes must respect the container limit (e.g. `-XX:MaxRAMPercentage=75` for JVMs, S7).
- Enforcement: Rancher container defaults (LimitRange) fill gaps; a policy (Kyverno/Gatekeeper, Q5) or the LimitRange's `maxLimitRequestRatio` can enforce the memory ratio per environment.



C2 Replicas

At least 2 replicas in prod, each able to carry the load alone.

Normal: load shared



drain / failure

One pod gone: the other carries it



Size requests so N-1 replicas carry the full load.


- **Prod: at least 2 replicas** for every stateless service; a single replica means downtime at every drain, upgrade and node failure.
- **Size for N-1 replicas:** with 2 replicas, one pod must carry the full load while the other restarts. With 3+, each pod can be smaller.
- **Quorum-based systems** (etcd, ZooKeeper, Kafka controllers, consensus DBs): 3 or 5 replicas, odd numbers.
- **Many small replicas beat few large ones:** they spread better, reschedule faster after a node failure and fit into the N+1 reserve. A single pod larger than the free space on any node can't be rescheduled at all.
- Quota must cover **replicas × requests + surge** (`maxSurge` pods during a rollout, ~10–25 %), or use `maxSurge: 0 / maxUnavailable: 1` for large pods (allocation model).
- **Single instance** workloads (stateful, licensed per instance, no leader election) are allowed but documented: they accept downtime on drains, and N+1 is checked for the pod itself (planner *Single instance* option).
- Non-prod: 1 replica is fine unless the environment is used to test HA behaviour (acc).




C3 Workload HA

Spread over nodes, a PDB that allows one disruption, real probes.

✗ Two replicas on one node = one replica



✓ Spread: `topologySpreadConstraints`



- **Spread replicas over failure domains:** `topologySpreadConstraints` on `kubernetes.io/hostname` (and on `topology.kubernetes.io/zone / rack` where nodes are labelled). Two replicas on one node are one replica.
- Prefer `topologySpreadConstraints` or **soft** (preferred) anti-affinity; hard anti-affinity with more replicas than nodes leaves pods Pending.
- **PodDisruptionBudget** for every prod workload with 2+ replicas: `maxUnavailable: 1` (or `minAvailable: N-1`).
- Never a PDB that allows zero disruptions (`minAvailable = replicas`, or `minAvailable: 1` with 1 replica): it blocks node drains and therefore every upgrade (E1).
- **Readiness and liveness probes** are part of HA: without readiness, traffic goes to pods that aren't ready; a too strict liveness probe restarts healthy pods under load.
- `terminationGracePeriodSeconds` and a `preStop` delay so in-flight requests finish during rollouts and drains.
- **Across clusters:** for workloads that need more than node-level HA, run them active/active or active/passive on two clusters; each cluster then needs quota for its share (or for the full load, if one cluster must carry it alone).
- Stateful HA (databases, queues) follows the operator's/vendor's topology; its storage must not be tied to a single node (local volumes don't move after a node failure).



D. Platform reliability

what we run · Is the platform itself as available as what we promise?

Platform team owns end to end. Tenants' HA is only as good as the platform underneath.


PLATFORM TEAM
Owns end to end

TEAMS
–



D1 Platform HA

3 control plane nodes, 3-node Rancher, 2+ ingress, node VMs on separate hosts.



- **Control plane:** 3 control plane / etcd nodes per cluster (odd number), separate from worker capacity.
- **Rancher local cluster** on 3 nodes; downstream clusters keep running if Rancher is down, but management and quota changes stop.
- **Ingress controller** with 2+ replicas spread over nodes, with its own PDB; it's in the path of every tenant.
- **Platform components** (monitoring, logging, DNS, CNI, storage drivers) run with platform priority classes and their requests are deducted from capacity (B1).
- **On-prem VMs:** node VMs of one cluster should sit on different hypervisor hosts (anti-affinity at the virtualisation layer), otherwise one host failure takes several nodes and N+1 is false (Q17).
- **Failure domains labelled:** nodes carry `topology.kubernetes.io/zone` (rack / host group) so workloads can spread over them (C3).
- **etcd snapshots** and Rancher backups stored off-cluster (E1).



E. Day 2 operations

keeping it true over time · Do the guarantees stay true after go-live?

Platform team runs it; teams take part in reviews and fix what's theirs.

PLATFORM TEAM

Runs maintenance, monitoring, reviews


TEAMS

Right-size, fix drift, join reviews



E1 Maintenance

Drains succeed unattended; upgrades go one node at a time.



- **Node drains** must succeed without manual work: PDBs allow disruption (C3), the N+1 reserve absorbs the moved pods (A1).
- **Upgrades** (Kubernetes, Rancher, OS, firmware) one node at a time; a drained node counts as the "+1". Don't start an upgrade on a cluster whose N+1 tile is red.
- **Backups**: etcd snapshots and Rancher backups stored off-cluster, restore tested at least yearly.
- **Housekeeping**: certificates and tokens rotated before expiry, unused namespaces and Projects removed (their quota returns to the pool), node sizes kept uniform when adding hardware.
- **Runbooks** for the recurring events: node failure, emergency quota increase, cluster at capacity, project onboarding and offboarding.



E2 Monitoring and alerts

Alert before tenants notice: N+1, quota at 90 %, Pending pods, OOMKills.

Alertmanager, or Kibana on the OTel metrics:



N+1 breached on a cluster (D5 in 07).



Project quota usage > 90 % for CPU or memory requests.



Pods Pending for more than a few minutes (quota exceeded or no node fits).



OOMKills and evictions per namespace.



CPU throttling above a threshold (CPU limits too tight).



Nodes NotReady or under memory/disk pressure.



E3 Reviews and changes

Monthly capacity check, quarterly project review, every quota change via Git.



- **Capacity reviews**: monthly look at allocation ratio and N+1 per cluster; quarterly review per project (quota vs. requests vs. usage), reclaim or right-size (process).
- **Quota changes** go through Git and review, like any other change (ADR-0001); emergency increases are allowed and followed up with a PR.
- **Right-sizing loop**: VPA/KRR recommendations → team adjusts requests → quota freed or increased → review.
- **Standards drift**: workloads without requests, memory limit ≠ request in prod, defaulted values — shown on the dashboard and Kibana, fixed by the owning team.



F. Finance

what it costs · What does capacity cost, and who pays for what?

Platform team owns the cost model and publishes the rates; finance approves them; projects pay from their budget.

PLATFORM TEAM

Owns the cost model, publishes unit rates


TEAMS

Pay for allocated quota from budget



F1 Cost model and unit rates

Projects pay for allocated quota at published rates; the reserve is priced in, not billed.



- **Monthly cost per node:** hardware amortisation (5 years), power and cooling, rack space and network, Rancher licence (F2), shared services (monitoring, logging, backup, registry) and the platform team's share.
- **Sellable capacity** = allocatable – platform components – N+1 reserve, and at most **80 %** of allocatable in prod.
- **Unit rates** = cluster cost ÷ sellable capacity, as **€ per vCPU-month** and **€ per GiB-month** (cost split between CPU and memory by a fixed ratio), published once a year.
- **The reserve is not free:** the N+1 node and the 20 % headroom are paid for through the rates. Nobody is billed for them separately, and nobody can buy them.
- **Projects pay for allocated quota**, not for usage (ADR-0002): budgets stay predictable, and giving back unused quota is the way to save.
- **Showback every month:** quota vs. requests vs. usage per project, on the same basis for every project.
- **Non-prod is cheaper per unit:** dev is overcommitted, so the same hardware sells more quota there.
- Shared and platform services are part of the rate unless finance decides to fund them separately (Q12).
- Unspent budget is fine; unused quota is not: it blocks capacity that others pay for.



F2 Rancher licence (CPU cores)

Rancher is licensed per CPU core: every core counts, sold or not; memory is licence-free.


- The Rancher subscription is counted on **CPU cores only**: memory, disk and the number of nodes don't change the licence cost.
- **Every core is licensed**, sold or not: the N+1 node and the 20 % reserve cost licence like the rest.
- **Licence cost goes into the CPU rate only**; the memory rate carries no licence.
- **Memory-rich nodes** (more GiB per core) give cheaper memory; don't buy cores the projects won't request.
- **Oversized CPU requests cost twice**: hardware and licence. CPU right-sizing lowers the licence need at renewal.
- **A new node is a licence step**: all its cores count from day one; plan node purchases with the licence renewal.
- To confirm in the contract: physical cores or threads, and whether control plane / etcd nodes count (Q16).



F3 Cluster size: minimum worker nodes and N+1

Prod: at least 3 worker nodes; from 5 equal nodes on, N+1 fits inside the 20 % reserve.

Share of a prod cluster that can be sold, by number of equal worker nodes



- The N+1 reserve is **one whole node** (hardware and licence): the fewer the nodes, the larger its share.
- Sellable share in prod = $\min((n - 1) \div n, 80 \%)$ of allocatable (after platform components): 2 nodes 50 %, 3 nodes 67 %, 4 nodes 75 %, **5 or more 80 %**.
- **Prod minimum: 3 worker nodes.** With 2, half the cluster is reserve, and a drain puts every workload on one node. 3 also lets quorum systems (3 replicas) spread over nodes.
- **Prod target: 5 or more equal nodes:** from there the 80 % ceiling, not N+1, is the limit, so the N+1 reserve costs nothing extra.
- **Acc/test:** 3 worker nodes when N+1 is kept (recommended), otherwise 2. **Dev:** at least 2, so a drain doesn't stop the environment.
- **Equal node sizes:** the largest node sets the reserve (A1); one big node makes the whole cluster more expensive.
- **Not too many tiny nodes** either: each node has fixed overhead (OS and kubelet reserves, DaemonSets for CNI, monitoring and logging, rack space, power). Prefer 5–10 equal nodes per prod cluster.
- **Control plane / etcd nodes** come on top (3 for HA, D1) and are not worker capacity.
- **Buy a node** when the next approved quota would take the cluster above 80 % or turn N+1 red.